

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

url = "C:/test.csv" # Update the path
data = pd.read_csv(url)
```

```
print(data.head())
print(data.info())
print(data.describe())
```

|                   | id | state | account_length | area_code     | international_plan |
|-------------------|----|-------|----------------|---------------|--------------------|
| voice_mail_plan \ |    |       |                |               |                    |
| 0                 | 1  | KS    | 128            | area_code_415 | no                 |
| yes               |    |       |                |               |                    |
| 1                 | 2  | AL    | 118            | area_code_510 | yes                |
| no                |    |       |                |               |                    |
| 2                 | 3  | IA    | 62             | area_code_415 | no                 |
| no                |    |       |                |               |                    |
| 3                 | 4  | VT    | 93             | area_code_510 | no                 |
| no                |    |       |                |               |                    |
| 4                 | 5  | NE    | 174            | area_code_415 | no                 |
| no                |    |       |                |               |                    |

|   | number_vmail_messages | total_day_minutes | total_day_calls | \ |
|---|-----------------------|-------------------|-----------------|---|
| 0 | 25                    | 265.1             | 110             |   |
| 1 | 0                     | 223.4             | 98              |   |
| 2 | 0                     | 120.7             | 70              |   |
| 3 | 0                     | 190.7             | 114             |   |
| 4 | 0                     | 124.3             | 76              |   |

|                    | total_day_charge | total_eve_minutes | total_eve_calls |
|--------------------|------------------|-------------------|-----------------|
| total_eve_charge \ |                  |                   |                 |
| 0                  | 45.07            | 197.4             | 99              |
| 16.78              |                  |                   |                 |
| 1                  | 37.98            | 220.6             | 101             |
| 18.75              |                  |                   |                 |
| 2                  | 20.52            | 307.2             | 76              |
| 26.11              |                  |                   |                 |
| 3                  | 32.42            | 218.2             | 111             |
| 18.55              |                  |                   |                 |
| 4                  | 21.13            | 277.1             | 112             |
| 23.55              |                  |                   |                 |

|   | total_night_minutes | total_night_calls | total_night_charge | \ |
|---|---------------------|-------------------|--------------------|---|
| 0 | 244.7               | 91                | 11.01              |   |
| 1 | 203.9               | 118               | 9.18               |   |
| 2 | 203.0               | 99                | 9.14               |   |
| 3 | 129.6               | 121               | 5.83               |   |

|   |                    |                  |                     |
|---|--------------------|------------------|---------------------|
| 4 | 250.7              | 115              | 11.28               |
|   | total_intl_minutes | total_intl_calls | total_intl_charge \ |
| 0 | 10.0               | 3                | 2.70                |
| 1 | 6.3                | 6                | 1.70                |
| 2 | 13.1               | 6                | 3.54                |
| 3 | 8.1                | 3                | 2.19                |
| 4 | 15.5               | 5                | 4.19                |

|   |                               |
|---|-------------------------------|
|   | number_customer_service_calls |
| 0 | 1                             |
| 1 | 0                             |
| 2 | 4                             |
| 3 | 3                             |
| 4 | 3                             |

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 750 entries, 0 to 749

Data columns (total 20 columns):

| #  | Column                        | Non-Null Count | Dtype   |
|----|-------------------------------|----------------|---------|
| 0  | id                            | 750 non-null   | int64   |
| 1  | state                         | 750 non-null   | object  |
| 2  | account_length                | 750 non-null   | int64   |
| 3  | area_code                     | 750 non-null   | object  |
| 4  | international_plan            | 750 non-null   | object  |
| 5  | voice_mail_plan               | 750 non-null   | object  |
| 6  | number_vmail_messages         | 750 non-null   | int64   |
| 7  | total_day_minutes             | 750 non-null   | float64 |
| 8  | total_day_calls               | 750 non-null   | int64   |
| 9  | total_day_charge              | 750 non-null   | float64 |
| 10 | total_eve_minutes             | 750 non-null   | float64 |
| 11 | total_eve_calls               | 750 non-null   | int64   |
| 12 | total_eve_charge              | 750 non-null   | float64 |
| 13 | total_night_minutes           | 750 non-null   | float64 |
| 14 | total_night_calls             | 750 non-null   | int64   |
| 15 | total_night_charge            | 750 non-null   | float64 |
| 16 | total_intl_minutes            | 750 non-null   | float64 |
| 17 | total_intl_calls              | 750 non-null   | int64   |
| 18 | total_intl_charge             | 750 non-null   | float64 |
| 19 | number_customer_service_calls | 750 non-null   | int64   |

dtypes: float64(8), int64(8), object(4)

memory usage: 117.3+ KB

None

|                     |           |                |                       |
|---------------------|-----------|----------------|-----------------------|
|                     | id        | account_length | number_vmail_messages |
| total_day_minutes \ |           |                |                       |
| count               | 750.00000 | 750.000000     | 750.000000            |
| 750.000000          |           |                |                       |
| mean                | 375.50000 | 100.385333     | 8.454667              |
| 180.454933          |           |                |                       |

|            |           |            |           |
|------------|-----------|------------|-----------|
| std        | 216.65064 | 39.699029  | 14.123712 |
| 53.258337  |           |            |           |
| min        | 1.00000   | 1.000000   | 0.000000  |
| 12.500000  |           |            |           |
| 25%        | 188.25000 | 74.000000  | 0.000000  |
| 146.625000 |           |            |           |
| 50%        | 375.50000 | 101.000000 | 0.000000  |
| 178.200000 |           |            |           |
| 75%        | 562.75000 | 126.000000 | 21.000000 |
| 215.975000 |           |            |           |
| max        | 750.00000 | 238.000000 | 51.000000 |
| 350.800000 |           |            |           |

|                   |                 |                  |                   |
|-------------------|-----------------|------------------|-------------------|
|                   | total_day_calls | total_day_charge | total_eve_minutes |
| total_eve_calls \ |                 |                  |                   |
| count             | 750.000000      | 750.000000       | 750.000000        |
| 750.000000        |                 |                  |                   |
| mean              | 100.721333      | 30.677920        | 203.258267        |
| 100.273333        |                 |                  |                   |
| std               | 19.718539       | 9.053756         | 52.185471         |
| 19.367535         |                 |                  |                   |
| min               | 39.000000       | 2.130000         | 31.200000         |
| 37.000000         |                 |                  |                   |
| 25%               | 88.000000       | 24.925000        | 166.800000        |
| 87.000000         |                 |                  |                   |
| 50%               | 101.000000      | 30.295000        | 203.350000        |
| 101.000000        |                 |                  |                   |
| 75%               | 114.000000      | 36.715000        | 235.975000        |
| 113.000000        |                 |                  |                   |
| max               | 163.000000      | 59.640000        | 363.700000        |
| 164.000000        |                 |                  |                   |

|       |                  |                     |                     |
|-------|------------------|---------------------|---------------------|
|       | total_eve_charge | total_night_minutes | total_night_calls \ |
| count | 750.000000       | 750.000000          | 750.000000          |
| mean  | 17.277080        | 199.619467          | 100.370667          |
| std   | 4.435638         | 51.531351           | 19.185238           |
| min   | 2.650000         | 50.900000           | 12.000000           |
| 25%   | 14.177500        | 164.475000          | 88.000000           |
| 50%   | 17.285000        | 199.450000          | 100.500000          |
| 75%   | 20.057500        | 234.800000          | 113.000000          |
| max   | 30.910000        | 364.300000          | 168.000000          |

|       |                    |                    |                    |
|-------|--------------------|--------------------|--------------------|
|       | total_night_charge | total_intl_minutes | total_intl_calls \ |
| count | 750.000000         | 750.000000         | 750.000000         |
| mean  | 8.982827           | 10.294133          | 4.485333           |
| std   | 2.318920           | 2.770340           | 2.421901           |
| min   | 2.290000           | 0.000000           | 0.000000           |
| 25%   | 7.402500           | 8.525000           | 3.000000           |
| 50%   | 8.975000           | 10.300000          | 4.000000           |
| 75%   | 10.565000          | 12.100000          | 6.000000           |

|     |           |           |           |
|-----|-----------|-----------|-----------|
| max | 16.390000 | 18.900000 | 19.000000 |
|-----|-----------|-----------|-----------|


|       |                   |                               |
|-------|-------------------|-------------------------------|
|       | total_intl_charge | number_customer_service_calls |
| count | 750.000000        | 750.000000                    |
| mean  | 2.779933          | 1.634667                      |
| std   | 0.747704          | 1.276207                      |
| min   | 0.000000          | 0.000000                      |
| 25%   | 2.305000          | 1.000000                      |
| 50%   | 2.780000          | 1.000000                      |
| 75%   | 3.270000          | 2.000000                      |
| max   | 5.100000          | 6.000000                      |


```
# Pairplot for general overview
sns.pairplot(data, hue='total_day_minutes') # Assuming
'international_plan' is the target column

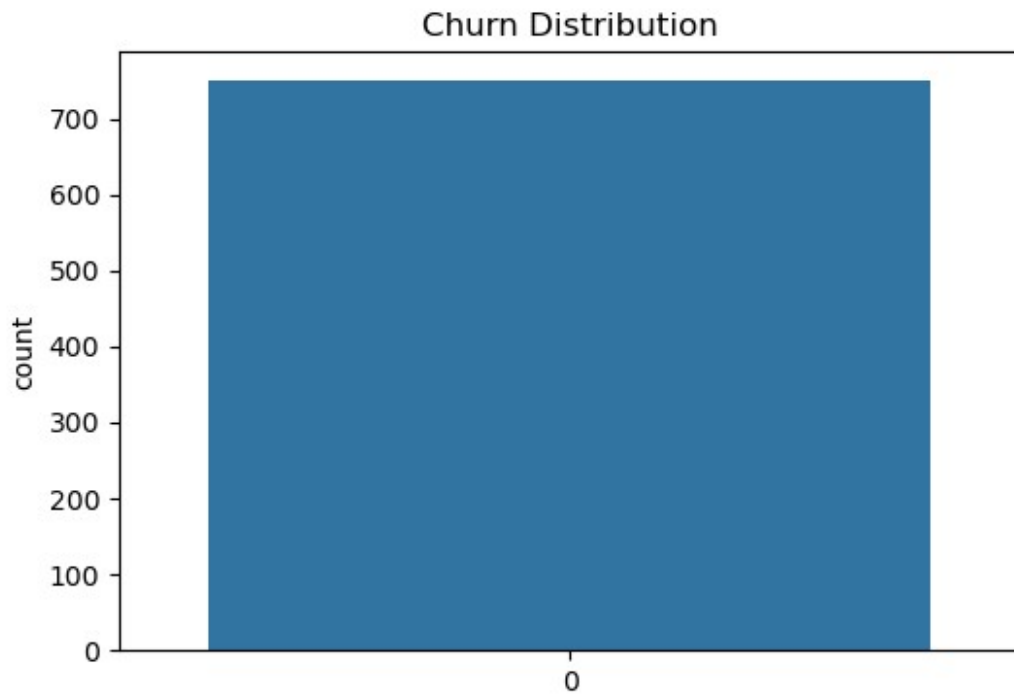
# Correlation heatmap
correlation_matrix = data.corr()
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()

# Distribution of Churn
plt.figure(figsize=(6, 4))
sns.countplot(data['total_day_minutes']) # Assuming
'international_plan' is the churn indicator
plt.title("Churn Distribution")
plt.show()
```



```
C:\Users\nishi\AppData\Local\Temp\ipykernel_2948\92574793.py:5:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only
valid columns or specify the value of numeric_only to silence this
warning.
    correlation_matrix = data.corr()
```








```
numerical_features = ['total_day_calls', 'total_day_charge',  
                      'total_eve_minutes'] # Update with your features  
for feature in numerical_features:  
    plt.figure(figsize=(6, 4))  
    sns.histplot(data[feature], kde=True)  
    plt.title(f"Distribution of {feature}")  
    plt.show()
```




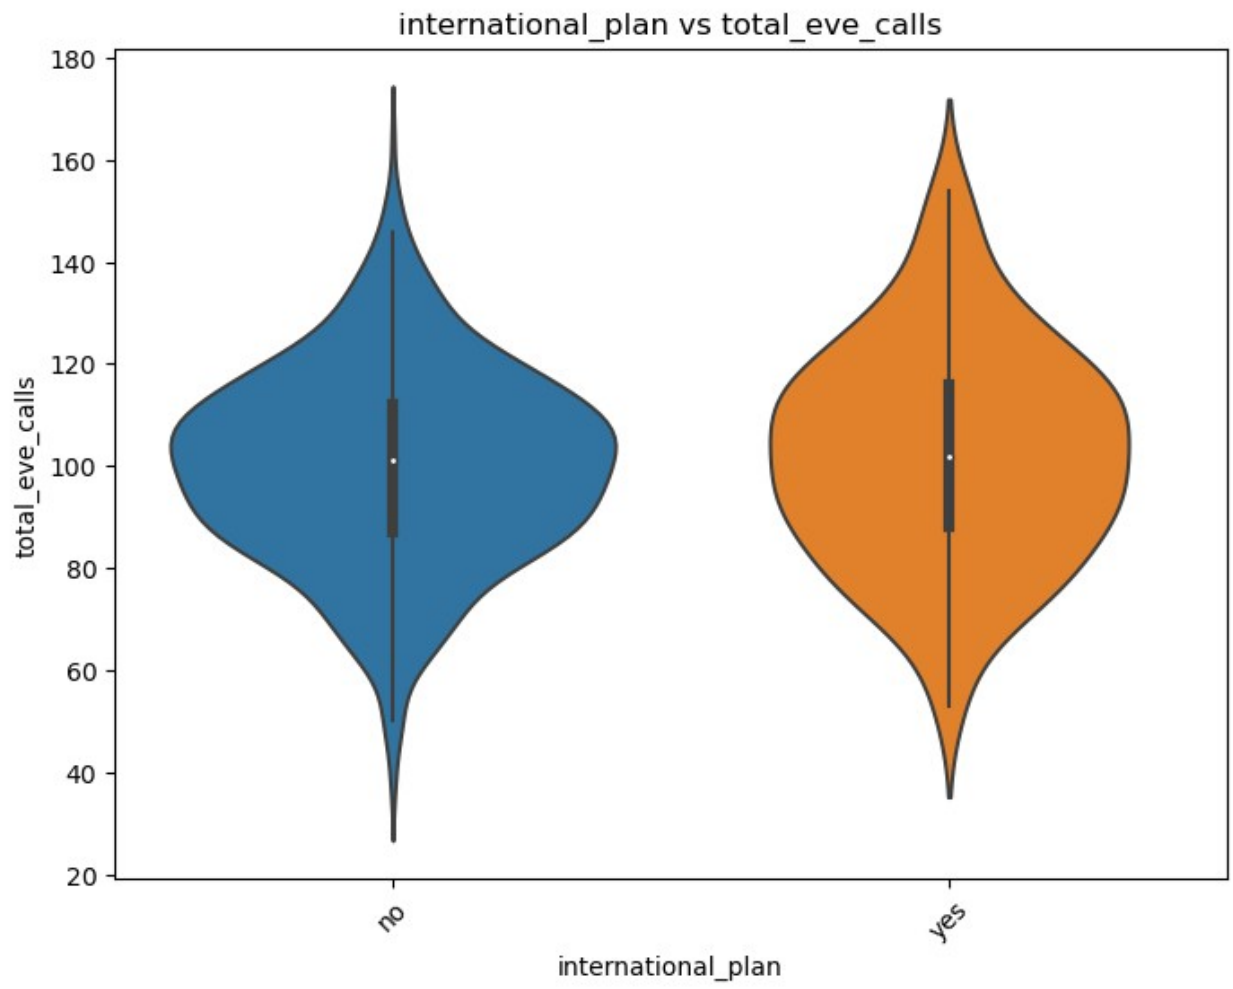


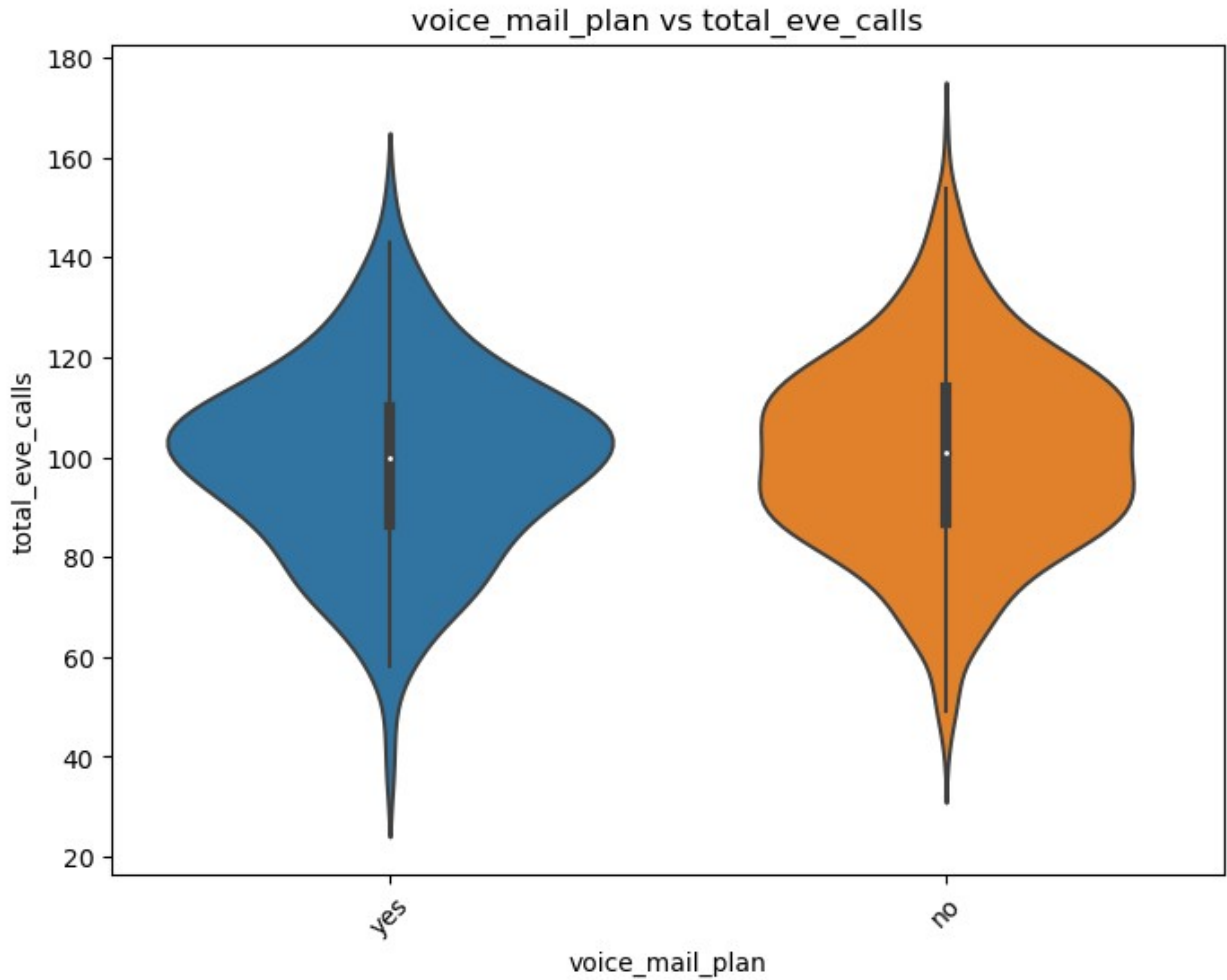


```
categorical_features = ['state', 'international_plan',  
                        'voice_mail_plan'] # Update with your features  
for feature in categorical_features:  
    plt.figure(figsize=(8, 6))  
    sns.violinplot(x=feature, y='total_eve_calls', data=data)  
    plt.title(f"{feature} vs total_eve_calls")  
    plt.xticks(rotation=45)  
    plt.show()
```

### state vs total\_eve\_calls







```
print(data.head())
```

|                   | id | state | account_length | area_code     | international_plan |
|-------------------|----|-------|----------------|---------------|--------------------|
| voice_mail_plan \ |    |       |                |               |                    |
| 0                 | 1  | KS    | 128            | area_code_415 | no                 |
| yes               |    |       |                |               |                    |
| 1                 | 2  | AL    | 118            | area_code_510 | yes                |
| no                |    |       |                |               |                    |
| 2                 | 3  | IA    | 62             | area_code_415 | no                 |
| no                |    |       |                |               |                    |
| 3                 | 4  | VT    | 93             | area_code_510 | no                 |
| no                |    |       |                |               |                    |
| 4                 | 5  | NE    | 174            | area_code_415 | no                 |
| no                |    |       |                |               |                    |

|   | number_vmail_messages | total_day_minutes | total_day_calls |
|---|-----------------------|-------------------|-----------------|
| 0 | 25                    | 265.1             | 110             |
| 1 | 0                     | 223.4             | 98              |
| 2 | 0                     | 120.7             | 70              |
| 3 | 0                     | 190.7             | 114             |

|   |   |       |    |
|---|---|-------|----|
| 4 | 0 | 124.3 | 76 |
|---|---|-------|----|

|                    | total_day_charge | total_eve_minutes | total_eve_calls |
|--------------------|------------------|-------------------|-----------------|
| total_eve_charge \ |                  |                   |                 |
| 0                  | 45.07            | 197.4             | 99              |
| 16.78              |                  |                   |                 |
| 1                  | 37.98            | 220.6             | 101             |
| 18.75              |                  |                   |                 |
| 2                  | 20.52            | 307.2             | 76              |
| 26.11              |                  |                   |                 |
| 3                  | 32.42            | 218.2             | 111             |
| 18.55              |                  |                   |                 |
| 4                  | 21.13            | 277.1             | 112             |
| 23.55              |                  |                   |                 |

|   | total_night_minutes | total_night_calls | total_night_charge \ |
|---|---------------------|-------------------|----------------------|
| 0 | 244.7               | 91                | 11.01                |
| 1 | 203.9               | 118               | 9.18                 |
| 2 | 203.0               | 99                | 9.14                 |
| 3 | 129.6               | 121               | 5.83                 |
| 4 | 250.7               | 115               | 11.28                |

|   | total_intl_minutes | total_intl_calls | total_intl_charge \ |
|---|--------------------|------------------|---------------------|
| 0 | 10.0               | 3                | 2.70                |
| 1 | 6.3                | 6                | 1.70                |
| 2 | 13.1               | 6                | 3.54                |
| 3 | 8.1                | 3                | 2.19                |
| 4 | 15.5               | 5                | 4.19                |

|   | number_customer_service_calls |
|---|-------------------------------|
| 0 | 1                             |
| 1 | 0                             |
| 2 | 4                             |
| 3 | 3                             |
| 4 | 3                             |

```
import pandas as pd

# Assuming you have a DataFrame named 'df' and want to delete columns
# 'column1' and 'column2'
columns_to_delete = ['area_code', 'state']

# Use the drop method to delete the specified columns
data.drop(columns=columns_to_delete, inplace=True)

# The 'inplace=True' argument modifies the DataFrame in place, so you
# don't need to reassign it.

print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 750 entries, 0 to 749
```

```
Data columns (total 16 columns):
```

| #  | Column                        | Non-Null Count | Dtype   |
|----|-------------------------------|----------------|---------|
| 0  | id                            | 750 non-null   | int64   |
| 1  | account_length                | 750 non-null   | int64   |
| 2  | number_vmail_messages         | 750 non-null   | int64   |
| 3  | total_day_minutes             | 750 non-null   | float64 |
| 4  | total_day_calls               | 750 non-null   | int64   |
| 5  | total_day_charge              | 750 non-null   | float64 |
| 6  | total_eve_minutes             | 750 non-null   | float64 |
| 7  | total_eve_calls               | 750 non-null   | int64   |
| 8  | total_eve_charge              | 750 non-null   | float64 |
| 9  | total_night_minutes           | 750 non-null   | float64 |
| 10 | total_night_calls             | 750 non-null   | int64   |
| 11 | total_night_charge            | 750 non-null   | float64 |
| 12 | total_intl_minutes            | 750 non-null   | float64 |
| 13 | total_intl_calls              | 750 non-null   | int64   |
| 14 | total_intl_charge             | 750 non-null   | float64 |
| 15 | number_customer_service_calls | 750 non-null   | int64   |

```
dtypes: float64(8), int64(8)
```

```
memory usage: 93.9 KB
```

```
None
```

```
print(data.isnull().sum())
```

```
# Depending on the columns with missing values, you can use methods  
like imputation.
```

|                               |   |
|-------------------------------|---|
| id                            | 0 |
| state                         | 0 |
| account_length                | 0 |
| area_code                     | 0 |
| number_vmail_messages         | 0 |
| total_day_minutes             | 0 |
| total_day_calls               | 0 |
| total_day_charge              | 0 |
| total_eve_minutes             | 0 |
| total_eve_calls               | 0 |
| total_eve_charge              | 0 |
| total_night_minutes           | 0 |
| total_night_calls             | 0 |
| total_night_charge            | 0 |
| total_intl_minutes            | 0 |
| total_intl_calls              | 0 |
| total_intl_charge             | 0 |
| number_customer_service_calls | 0 |

```
dtype: int64
```

```

from sklearn.model_selection import train_test_split

X = data.drop(columns=["total_intl_calls"])
y = data["total_intl_calls"]

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

print(X.shape, X_train.shape, X_test.shape)

(750, 15) (600, 15) (150, 15)

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train[numerical_features] =
scaler.fit_transform(X_train[numerical_features])
X_test[numerical_features] =
scaler.transform(X_test[numerical_features])

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
model = LogisticRegression(max_iter=1000)# Import LogisticRegression
#from sklearn.tree import DecisionTreeClassifier
#from sklearn.ensemble import RandomForestClassifier
#from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
# Preprocess the categorical variables using one-hot encoding
#data = pd.get_dummies(data, columns=["Geography", "Gender"])

# Split into features (X) and target (y)
X = data.drop("number_customer_service_calls", axis=1)
y = data["number_customer_service_calls"]

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define the models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    #"Decision Tree": DecisionTreeClassifier(),
    #"Random Forest": RandomForestClassifier(),
    #"Support Vector Machine": SVC()
}

results = {}

for name, model in models.items():
    model.fit(X_train, y_train)

```

```

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted',
zero_division=0) # Specify 'weighted' for multiclass
recall = recall_score(y_test, y_pred, average='weighted') #
Specify 'weighted' for multiclass
f1 = f1_score(y_test, y_pred, average='weighted') # Specify
'weighted' for multiclass

results[name] = {"Accuracy": accuracy, "Precision": precision,
"Recall": recall, "F1": f1}

# Display evaluation results
print("Model Evaluation Results:")
for name, metrics in results.items():
    print(f"Model: {name}")
    print(f"Accuracy: {metrics['Accuracy']:.2f}")
    print(f"Precision: {metrics['Precision']:.2f}")
    print(f"Recall: {metrics['Recall']:.2f}")
    print(f"F1 Score: {metrics['F1']:.2f}")
    print()

```

```

Model Evaluation Results:
Model: Logistic Regression
Accuracy: 0.34
Precision: 0.17
Recall: 0.34
F1 Score: 0.21

```

```

C:\Users\nishi\anaconda3\Lib\site-packages\sklearn\linear_model\
_logistic.py:460: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(

print(data.head())

```

|   | id | account_length | number_vmail_messages | total_day_minutes | \ |
|---|----|----------------|-----------------------|-------------------|---|
| 0 | 1  | 128            | 25                    | 265.1             |   |
| 1 | 2  | 118            | 0                     | 223.4             |   |
| 2 | 3  | 62             | 0                     | 120.7             |   |



|   |   |     |   |       |
|---|---|-----|---|-------|
| 3 | 4 | 93  | 0 | 190.7 |
| 4 | 5 | 174 | 0 | 124.3 |

|                   | total_day_calls | total_day_charge | total_eve_minutes |
|-------------------|-----------------|------------------|-------------------|
| total_eve_calls \ |                 |                  |                   |
| 0                 | 110             | 45.07            | 197.4             |
| 99                |                 |                  |                   |
| 1                 | 98              | 37.98            | 220.6             |
| 101               |                 |                  |                   |
| 2                 | 70              | 20.52            | 307.2             |
| 76                |                 |                  |                   |
| 3                 | 114             | 32.42            | 218.2             |
| 111               |                 |                  |                   |
| 4                 | 76              | 21.13            | 277.1             |
| 112               |                 |                  |                   |

|   | total_eve_charge | total_night_minutes | total_night_calls | \ |
|---|------------------|---------------------|-------------------|---|
| 0 | 16.78            | 244.7               | 91                |   |
| 1 | 18.75            | 203.9               | 118               |   |
| 2 | 26.11            | 203.0               | 99                |   |
| 3 | 18.55            | 129.6               | 121               |   |
| 4 | 23.55            | 250.7               | 115               |   |

|   | total_night_charge | total_intl_minutes | total_intl_calls | \ |
|---|--------------------|--------------------|------------------|---|
| 0 | 11.01              | 10.0               | 3                |   |
| 1 | 9.18               | 6.3                | 6                |   |
| 2 | 9.14               | 13.1               | 6                |   |
| 3 | 5.83               | 8.1                | 3                |   |
| 4 | 11.28              | 15.5               | 5                |   |

|   | total_intl_charge | number_customer_service_calls |
|---|-------------------|-------------------------------|
| 0 | 2.70              | 1                             |
| 1 | 1.70              | 0                             |
| 2 | 3.54              | 4                             |
| 3 | 2.19              | 3                             |
| 4 | 4.19              | 3                             |

```
best_model = max(results, key=lambda k: results[k]["F1"])
print("Best Performing Model:", best_model)
```

Best Performing Model: Logistic Regression

```
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

*# Input data*

```
input_data = (175, 258, 587, 123, 456, 258, 236, 147, 159, 368, 157,
147, 142, 124, 152)
```

*# Changing the input\_data to a numpy array*

```

input_data_as_numpy_array = np.asarray(input_data)

# Reshape the array as we are predicting for one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1, -1)

# Assuming you have a DataFrame named 'data' and want to predict
# 'number_customer_service_calls'
# Replace 'X' and 'y' with your actual features and target variable
X = data.drop("number_customer_service_calls", axis=1)
y = data["number_customer_service_calls"]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Create and train a Logistic Regression model# Create and train a
Logistic Regression model with increased max_iter
logistic_regression = LogisticRegression(max_iter=1000)
logistic_regression.fit(X_train, y_train)

# Now, you can make predictions with the trained model
prediction = logistic_regression.predict(input_data_reshaped)

if prediction[0] == 0:
    print("Churn: No")
else:
    print("Churn: Yes")

```

Churn: Yes

```

C:\Users\nishi\anaconda3\Lib\site-packages\sklearn\linear_model\
_logistic.py:460: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(

```

```

C:\Users\nishi\anaconda3\Lib\site-packages\sklearn\base.py:464:
UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
warnings.warn(

```